

ofproto/detrace – the missing link

Aleš Musil, Dumitru Ceară, 2024

The network in the kernel..

```
ufid:afa73c50-ee58-4f65-8cd2-7fb5d0271ae4,
recirc_id(0),dp_hash(0/0),skb_priority(0/0),in_port(client),skb_mark(0/0),ct_state(0/0x21),ct_zone(0/0),ct_mark(0/0),ct_label(0/0),eth(src=00:00:00:00:31:10,dst=00:00:00:00:31:00),eth_type(0x0800),ipv4(src=10.244.1.4,dst=10.244.1.42,proto=6,tos=0/0,ttl=64,frag=no),tcp(src=0/0,dst=0/0),tcp_flags(0/0), packets:61, bytes:4142, used:0.072s,
flags:P, dp:ovs, actions:ct(zone=2,nat),recirc(0x25)

ufid:2cc4a33c-c9fe-4cc0-aa37-adbc2494a856,
recirc_id(0x25),dp_hash(0/0),skb_priority(0/0),in_port(client),skb_mark(0/0),ct_state(0x22/0x27),ct_zone(0/0),ct_mark(0x2/0xe),ct_label(0/0),eth(src=00:00:00:00:31:10,dst=00:00:00:00:31:00),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=10.244.2.0/255.255.255.0,proto=6,tos=0/0,ttl=64,frag=no),tcp(src=0/0,dst=0/0),tcp_flags(0/0),
packets:60, bytes:4076, used:0.072s, flags:P, dp:ovs, actions:set(eth(src=00:00:00:00:10:00,dst=00:00:00:00:20:00)),set(ipv4(ttl=63)),ct(zone=2,nat),recirc(0x28)

ufid:15b43126-bae2-48eb-bc05-f7580b46ff61,
recirc_id(0x28),dp_hash(0/0),skb_priority(0/0),in_port(client),skb_mark(0/0),ct_state(0x20/0x21),ct_zone(0/0),ct_mark(0/0),ct_label(0/0),eth(src=00:00:00:00:10:00,dst=00:00:00:00:20:00),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0.0/0.0.0.0,proto=0/0,tos=0/0x3,ttl=0/0,frag=no), packets:60, bytes:4076, used:0.072s, flags:P, dp:ovs,
actions:ct_clear,set(tunnel(tun_id=0xa,src=170.168.0.4,dst=170.168.0.5,ttl=64,tp_dst=6081,geneve({class=0x102,type=0x80,len=4,0xa0014})),flags(df|csum|key))),genev_sys_6081

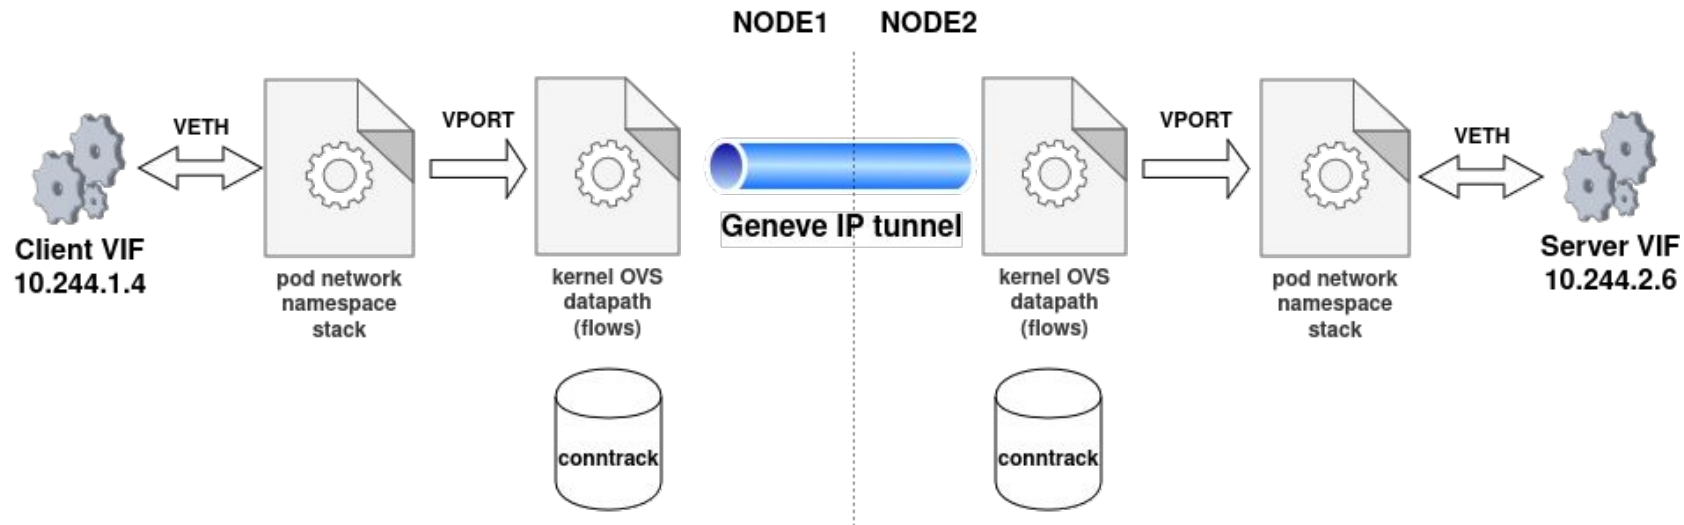
ufid:6350dfed-cee1-4feb-9570-9c8ae3c50b80,
recirc_id(0),dp_hash(0/0),skb_priority(0/0),tunnel(tun_id=0xa,src=170.168.0.5,dst=170.168.0.4,ttl=0/0,geneve({class=0x102,type=0x80,len=4,0x14000a/0x7fffffff})),flags(-df+csum+key)),in_port(genev_sys_6081),skb_mark(0/0),ct_state(0/0x27),ct_zone(0/0),ct_mark(0/0xe),ct_label(0/0),eth(src=00:00:00:00:20:00,dst=00:00:00:00:10:00),eth_type(0x0800),ipv4(src=10.244.2.0/255.255.254.0,dst=10.244.1.4,proto=6,tos=0/0,ttl=63,frag=no),tcp(src=0/0,dst=0/0),tcp_flags(0/0), packets:60, bytes:3962, used:0.072s, flags:P,
dp:ovs, actions:set(eth(src=00:00:00:00:31:00,dst=00:00:00:00:31:10)),set(ipv4(ttl=62)),ct(zone=2,nat),recirc(0x29)

ufid:81838d8b-4046-49b0-b20f-08f2bf7249b4,
recirc_id(0x29),dp_hash(0/0),skb_priority(0/0),tunnel(tun_id=0xa,src=170.168.0.5,dst=170.168.0.4,ttl=0/0,geneve({class=0/0,type=0/0,len=0/0}{class=0/0,type=0/0,len=0/0})),flags(-df+csum+key)),in_port(genev_sys_6081),skb_mark(0/0),ct_state(0x20/0x21),ct_zone(0/0),ct_mark(0/0),ct_label(0/0),eth(src=00:00:00:00:31:00,dst=00:00:00:00:31:10),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0.0/0.0.0.0,proto=0/0,tos=0/0,ttl=0/0,frag=no), packets:60, bytes:3962, used:0.072s, flags:P, dp:ovs,
actions:ct_clear,client
```

How do we understand and debug this?

- an overview of the topology we're trying to debug
- walk our way up the stack and connect the dots between abstraction layers
- take advantage of OVS/OVN tools already available

The “kernel view”



- virtual networking through OVS (kernel) datapath flow table:
 - **match:** (masked) packet fields, conntrack state / label / mark, ingress port, recirc_id (if multiple passes required - conntrack), etc.
 - **actions:** drop, output to port, send to conntrack (and recirculate), set packet fields or metadata, etc.
 - **UFID:** *unique flow identifier*
 - quickly identify flows between userspace (ovs-vswitchd) / kernel (OVS kmod)
 - mapping between userspace context (e.g., OpenFlow rules) and flows

The “kernel view”

```
$ ovs-appctl dpctl/dump-flows -m
```

uuid:afa73c50-ee58-4f65-8cd2-7fb5d0271ae4

match: recirc_id(0), in_port(client), eth(src=00:00:00:00:31:10, dst=00:00:00:00:31:00), ipv4(src=**10.244.1.4**, dst=**10.244.1.42**), packets:61

actions: ct(zone=2, nat), recirc(0x25)

ct_state=+est+trk

uuid:2cc4a33c-c9fe-4cc0-aa37-adbc2494a856


match: recirc_id(0x25), in_port(client), ct_state(0x22/0x27), ct_mark(0x2/0xe), eth(src=00:00:00:00:31:10, dst=00:00:00:00:31:00),
ipv4(dst=**10.244.2.0**/255.255.255.0)

actions: set(eth(src=00:00:00:00:10:00, dst=00:00:00:00:20:00)), set(ipv4(ttl=63)), ct(zone=2, nat), recirc(0x28)

- traffic flowing from the client's veth to the geneve tunnel
- conntrack lookup and a conntrack entry (established session) being matched
- packet destination IP changing (DNAT)
- packet MAC addresses changing and TTL decrementing (routing?)
- and more..

That's useful but can we figure out why that's happening?

The “OpenFlow view”



Rule format:

- table-id, priority, match, actions
- **cookie:** ovn-controller leaves “bread crumbs”, (IDs of OVN Southbound database records):
Logical_Flows, Port_Bindings, Load_Balancers, etc.

```
$ ovs-appctl list-commands | grep trace
```

```
ofproto/detrace      UFID  [pmd=PMD-ID]
```

```
ofproto/trace        {[dp_name] odp_flow | bridge br_flow} [OPTIONS...] [-generate|packet]
```

```
ofproto/trace-packet-out [-consistent] {[dp_name] odp_flow | bridge br_flow} [OPTIONS...] [-generate|packet] actions
```

```
$ grep -A1 ofproto/detrace NEWS
```

- * Added 'ofproto/detrace' command that outputs the set of OpenFlow rules and groups that contributed to the creation of a specific datapath flow.


UFID: *unique flow identifier*

- quickly identify flows between userspace (ovs-vswitchd) / kernel (OVS kmod)
- mapping between userspace context (e.g., OpenFlow rules) and flows

pmd/PMD-ID: *(DPDK) poll mode driver / poll mode driver ID*

- Core to which this queue should be pinned. OVS_CORE_UNSPEC if the queue doesn't need to be pinned to a particular core.
- for the kernel (netlink) datapath: PMD_ID_NULL (INT_MAX)

The "OVN view"



For each node:

- separate OVN "cluster" connected to a distributed **transit Logical_Switch** (geneve)
- **node Logical_Switch** (*ovn-worker-x*)
- **Logical_Router** (*ovn_cluster_router*) for E-W vs N-S routing decisions
- ...

For each pod, on the node where it runs:

- **Logical_Switch_Port** connected to the node logical switch

For each service, on each node:

- **Load_Balancer** attached to the node logical router

kernel -> OpenFlow -> OVN Southbound

OpenFlow:

\$ **ovs-appctl** ofproto/detrace **ufid:afa73c50-ee58-4f65-8cd2-7fb5d0271ae4**

cookie=**0xaed19b4d**,table_id=0,priority=100,**match**=in_port=4,
actions=set_field:0xb/0xffff->reg13,set_field:0x9->reg11,set_field:0x8->reg12,**set_field:0x2->metadata,set_field:0x2->reg14**,set_field:0/0xffff0000->reg13,resubmit(,8)

cookie=**0x1d826318**, table_id=14,priority=100,**match**=ip,metadata=0x1,nw_dst=**10.244.1.42**,actions=ct(table=15,zone=NXM_NX_REG11[0..15],nat)

OVN Southbound DB:

\$ **ovn-sbctl** list **Port_Binding** **aed19b4d**

_uuid : **aed19b4d**-7f17-4654-b092-9c0d534a0c0f
logical_port : **client**
mac : ["00:00:00:00:31:10 10.244.1.4"]

Corresponding to NB 'client'
logical port

\$ **ovn-sbctl** list **Logical_Flow** **1d826318**

external_ids : {source="northd.c:12397", stage-hint="**501b7c9c**", stage-name=lr_in_defrag}
priority : 100
match : "ip && ip4.dst == 10.244.1.42"
actions : "ct_dnat;"

A load balancer related
OpenFlow rule hit (ct_dnat)

OVN Southbound -> OVN Northbound

OVN Southbound DB:

\$ **ovn-sbctl** list **Port_Binding** aed19b4d

_uuid : aed19b4d-7f17-4654-b092-9c0d534a0c0f
logical_port : client
mac : ["0a:58:0a:f4:01:04 10.244.1.4"]

stage-hint by ovn-northd to
NB objects

\$ **ovn-sbctl** list **Logical_Flow** 1d826318

external_ids : {source="northd.c:12397", stage-hint="501b7c9c", stage-name=lr_in_defrag}
priority : 100
match : "ip && ip4.dst == 10.244.1.42"
actions : "ct_dnat;"

OVN Northbound DB:

\$ **ovn-nbctl** list **Logical_Switch_Port** client

addresses : ["00:00:00:00:31:10 10.244.1.4"]
name : client
up : true

\$ **ovn-nbctl** list **Load_Balancer** 501b7c9c

_uuid : 501b7c9c-b219-4159-b256-1fe633dafc15
name : "lb0"
protocol : tcp
vips : {"10.244.1.42:80"="10.244.2.6:80"}

Load balancer
VIP/backend mapping

\$ man ovn-detrace
ovn-detrace(1)

OVN Manual

ovn-detrace(1)

NAME

ovn-detrace - convert ``ovs-appctl ofproto/trace" (*) output to combine OVN logical flow information.

SYNOPSIS

ovn-detrace < file

DESCRIPTION

The ovn-detrace program reads ovs-appctl ofproto/trace output on stdin, looking for flow cookies, and expand each cookie with corresponding OVN logical flows. It expands logical flow further with the north-bound information e.g. the ACL that generated the logical flow, when relevant.

[...]

(*) works perfectly fine with *ovs-appctl ofproto/detrace* **output too.**

cookie=**0x1d826318**,table_id=14,priority=100,**match**=ip,metadata=0x1,nw_dst=**10.244.1.42**,
actions=ct(table=15,zone=NXM_NX_REG11[0..15],nat)

* Logical datapaths:

* "gw" (0ed8f608-8736-4158-80b1-c7341f7659a4) [ingress]

* Logical flow: table=6 (lr_in_defrag), priority=100, **match**=(ip && ip4.dst == **10.244.1.42**), **actions**=(ct_dnat;)

* Load Balancer: **lb0** protocol ['tcp'] vips {'**10.244.1.42:80**': '**10.244.2.6:80**'} ip_port_mappings {}

Connecting the dots

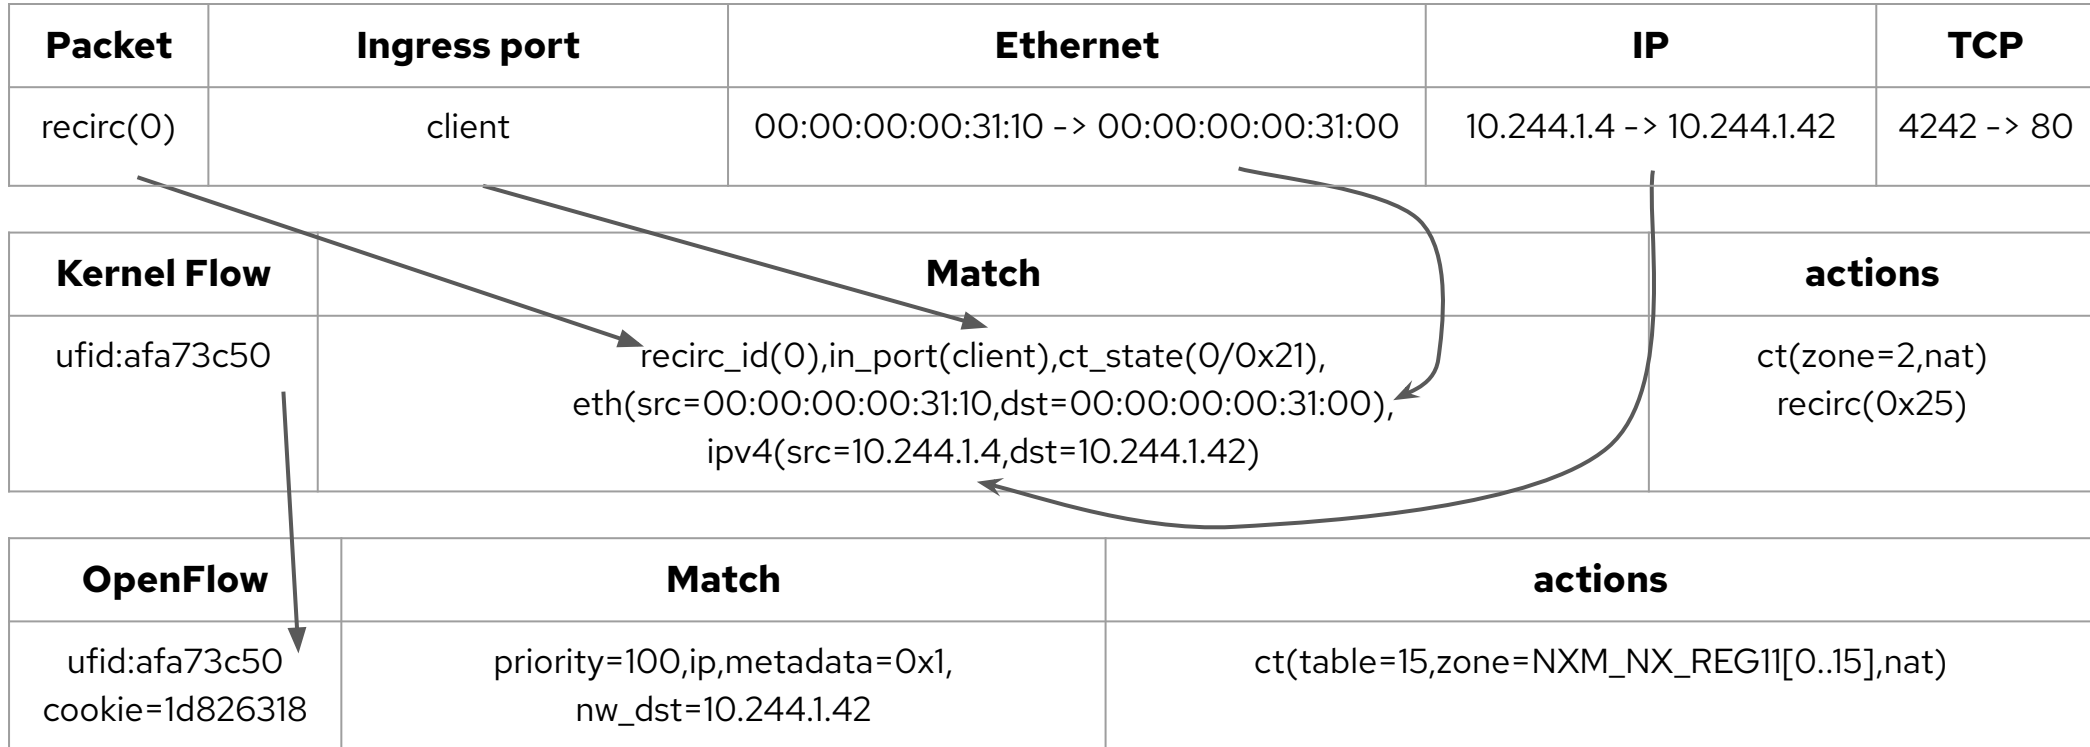
| Packet | Ingress port | Ethernet | IP | TCP |
|-----------|--------------|--|---------------------------|------------|
| recirc(0) | client | 00:00:00:00:31:10 -> 00:00:00:00:31:00 | 10.244.1.4 -> 10.244.1.42 | 4242 -> 80 |

Connecting the dots

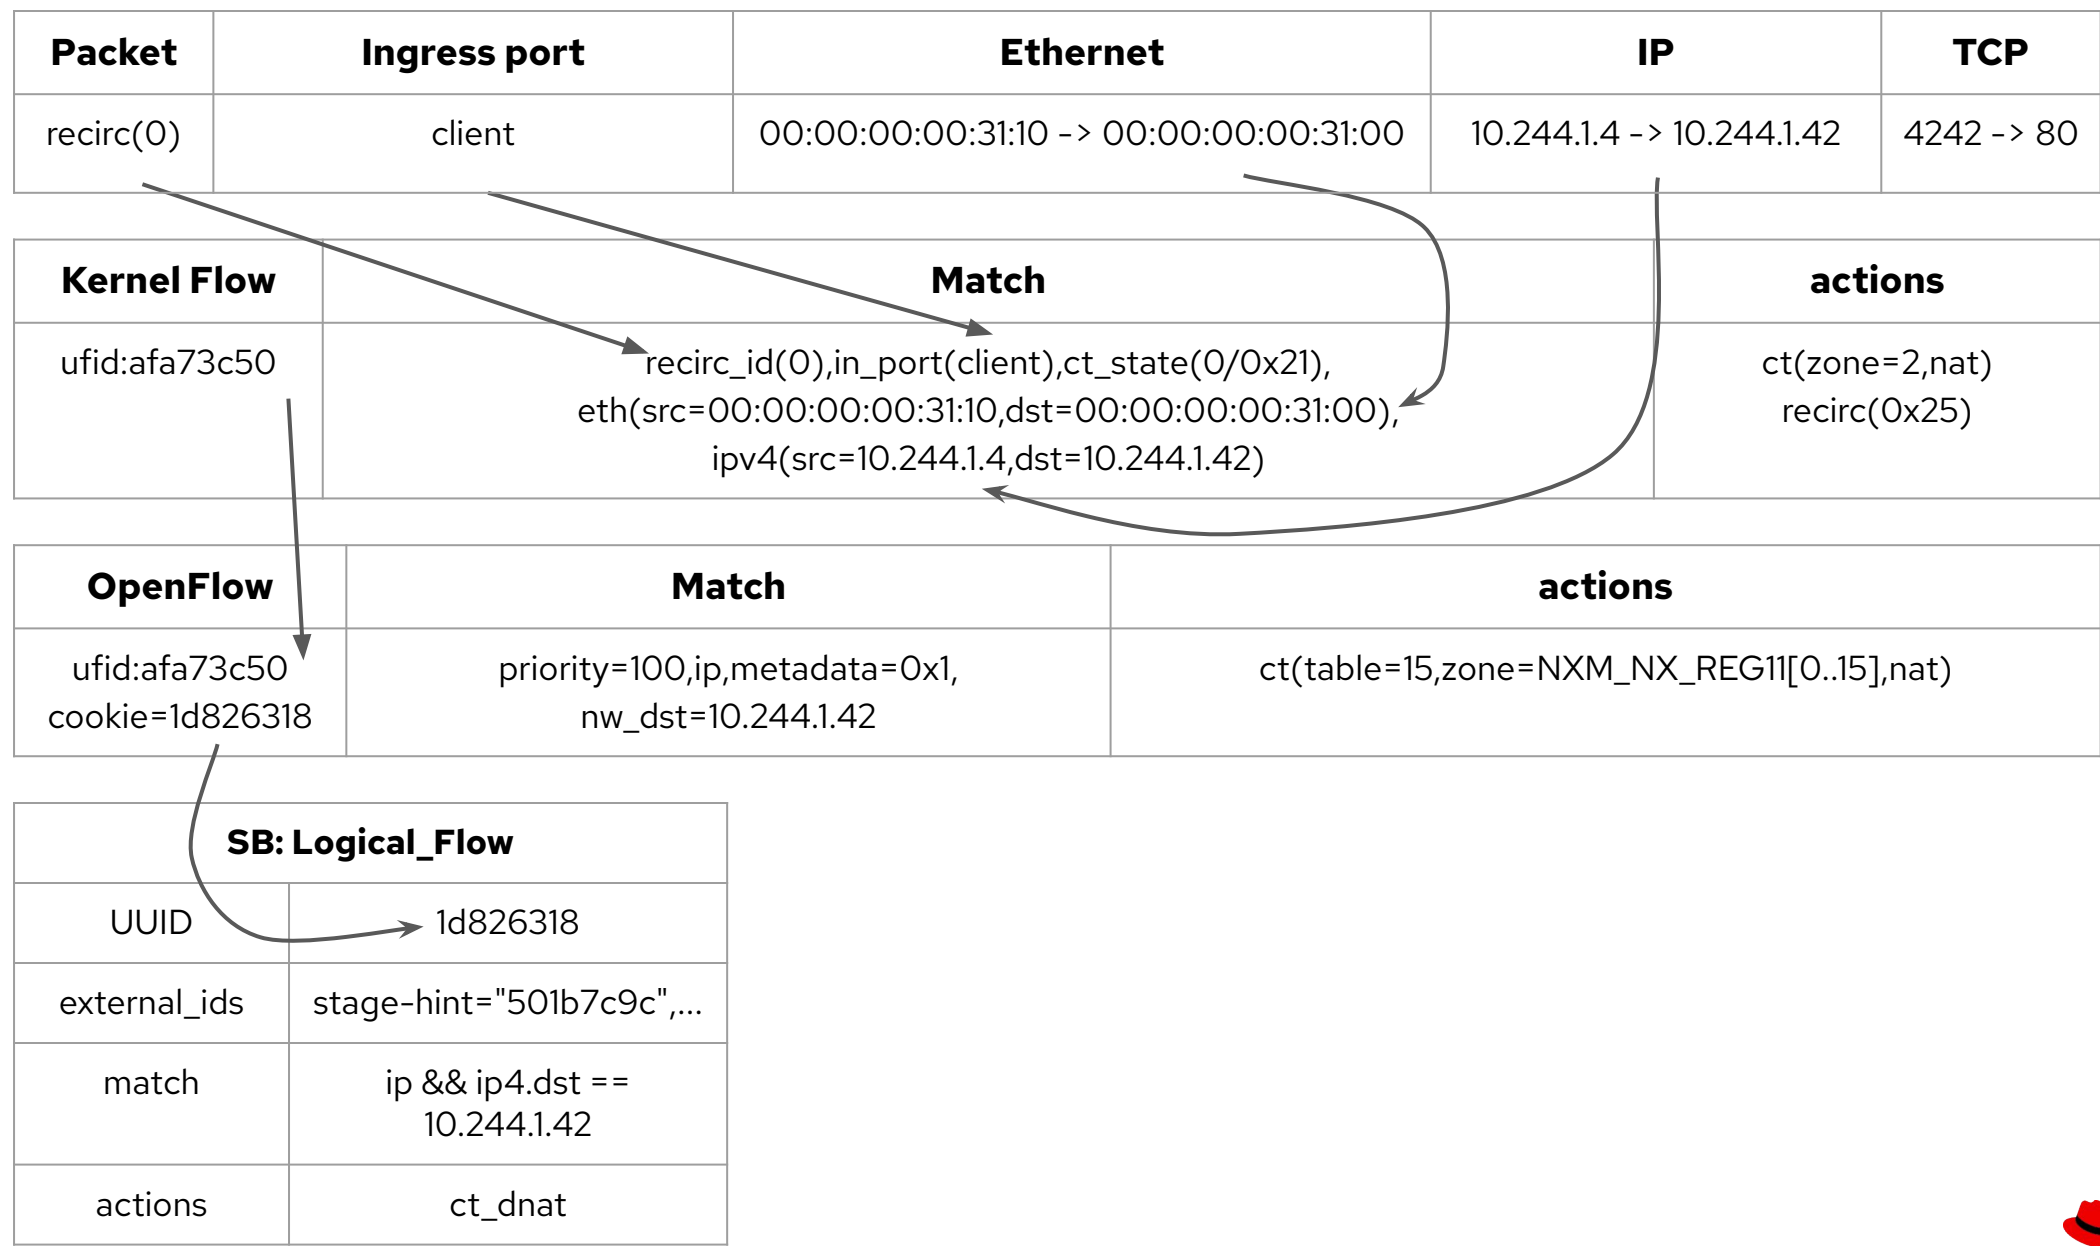
| Packet | Ingress port | Ethernet | IP | TCP |
|-----------|--------------|--|---------------------------|------------|
| recirc(0) | client | 00:00:00:00:31:10 -> 00:00:00:00:31:00 | 10.244.1.4 -> 10.244.1.42 | 4242 -> 80 |

| Kernel Flow | Match | actions |
|---------------|---|--------------------------------|
| ufid:afa73c50 | recirc_id(0),in_port(client),ct_state(0/0x21), eth(src=00:00:00:00:31:10,dst=00:00:00:00:31:00), ipv4(src=10.244.1.4,dst=10.244.1.42) | ct(zone=2,nat) recirc(0x25) |

Connecting the dots



Connecting the dots



Connecting the dots

| Packet | Ingress port | Ethernet | IP | TCP |
|-----------|--------------|--|---------------------------|------------|
| recirc(0) | client | 00:00:00:00:31:10 -> 00:00:00:00:31:00 | 10.244.1.4 -> 10.244.1.42 | 4242 -> 80 |

| Kernel Flow | Match | actions |
|---------------|---|--------------------------------|
| ufid:afa73c50 | recirc_id(0),in_port(client),ct_state(0/0x21), eth(src=00:00:00:00:31:10,dst=00:00:00:00:31:00), ipv4(src=10.244.1.4,dst=10.244.1.42) | ct(zone=2,nat) recirc(0x25) |

| OpenFlow | Match | actions |
|----------------------------------|---|---|
| ufid:afa73c50 cookie=1d826318 | priority=100,ip,metadata=0x1, nw_dst=10.244.1.42 | ct(table=15,zone=NXM_NX_REG11[0..15],nat) |

| SB: Logical_Flow | |
|------------------|---------------------------------|
| UUID | 1d826318 |
| external_ids | stage-hint="501b7c9c",... |
| match | ip && ip4.dst == 10.244.1.42 |
| actions | ct_dnat |

| NB: Load_Balancer | |
|-------------------|---------------------------------|
| UUID | 501b7c9c |
| name | lb0 |
| vips | 10.244.1.42:80 -> 10.244.2.6:80 |
| protocol | tcp |

Key takeaways?

- **there are no mysteries in networking**
 - it is complex and might look scary at first
 - **but** with the right tools it can become more **readable** and **easier** to understand
- **there's a link between abstractions at all the levels in the stack**
 - *from* highly optimized kernel datapath flows
 - *to* OVN (logical) routers and switches
- **use available tools** to better understand what the network is doing
- **when in doubt remember to check for:** UFIDs -> cookies -> stage-hints -> external-ids

Thank you!