

# (P)sampling kubernetes network policies

Dumitru Ceară

Nadia Pinaeva

Adrián Moreno


# Kubernetes

Kubernetes, also known as k8s, is a container orchestration engine.

It groups containers that make up an application into logical units called Pods.

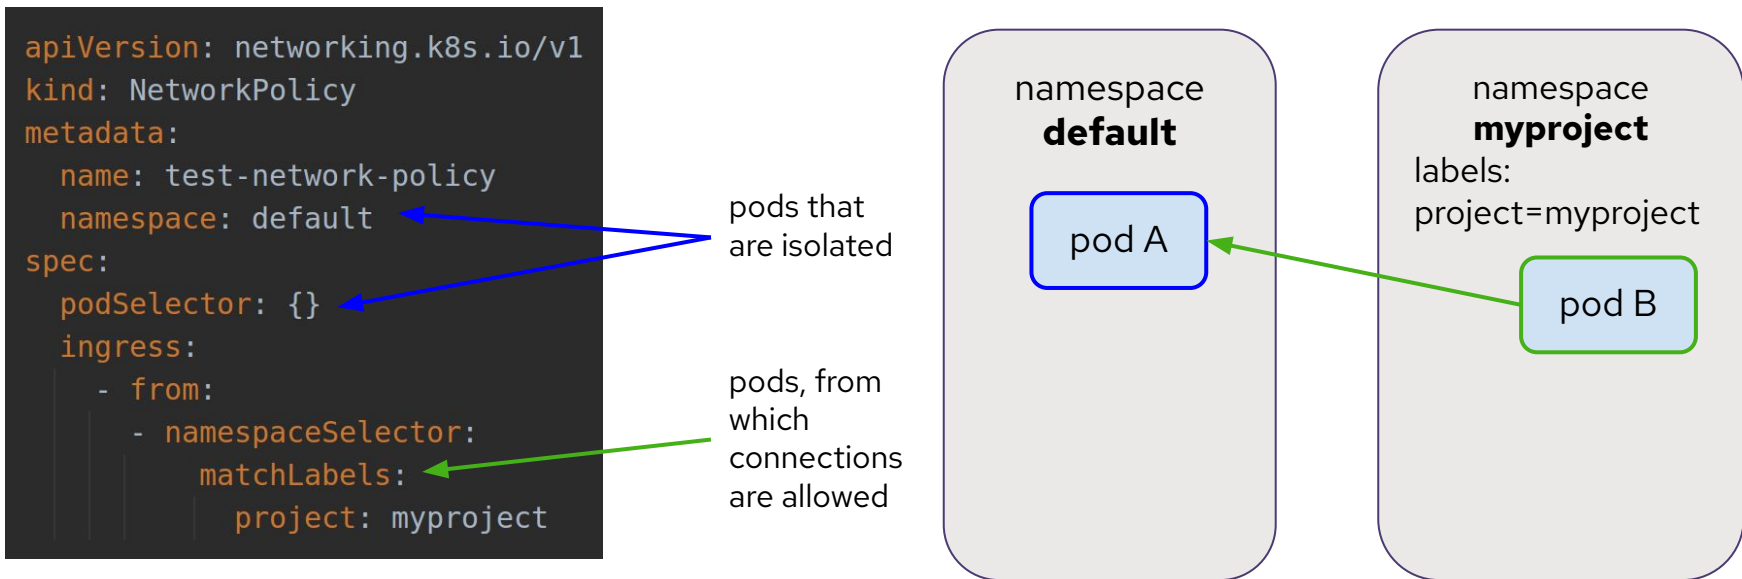
Cluster resources, including Pods, are isolated using namespaces.

Labels are key/value pairs that are attached to objects such as Pods and Namespaces. Labels can be used to organize and to select subsets of objects.



# Kubernetes Network Policy


NetworkPolicy is a kubernetes API that ensures network security, by specifying connections that should be allowed for a namespace. All the other connections for a namespace will be denied.



## Observability problem


In real clusters, the amount of Network Policies may go beyond thousands.

It gets difficult to understand, why exactly some connection is allowed or denied.

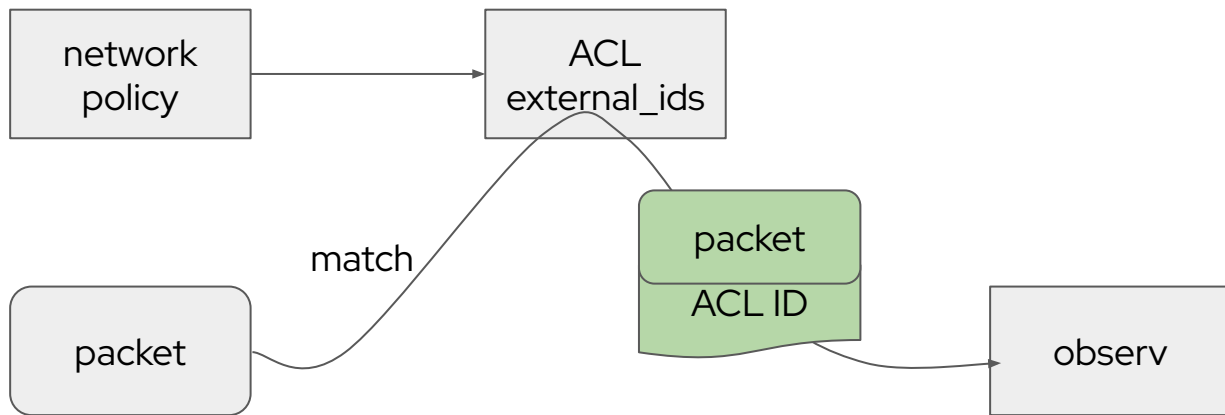


# OVN-Kubernetes

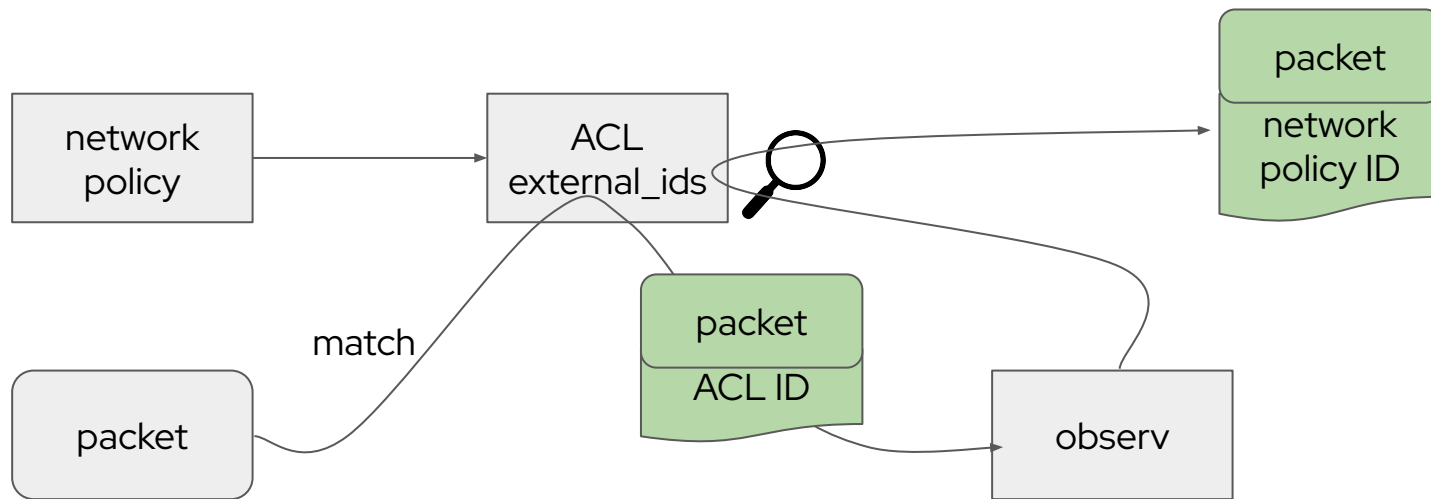
OVN-kubernetes uses OVN ACLs to implement k8s Network Policy.



OVN-kubernetes uses ACLs to implement k8s Network Policy.

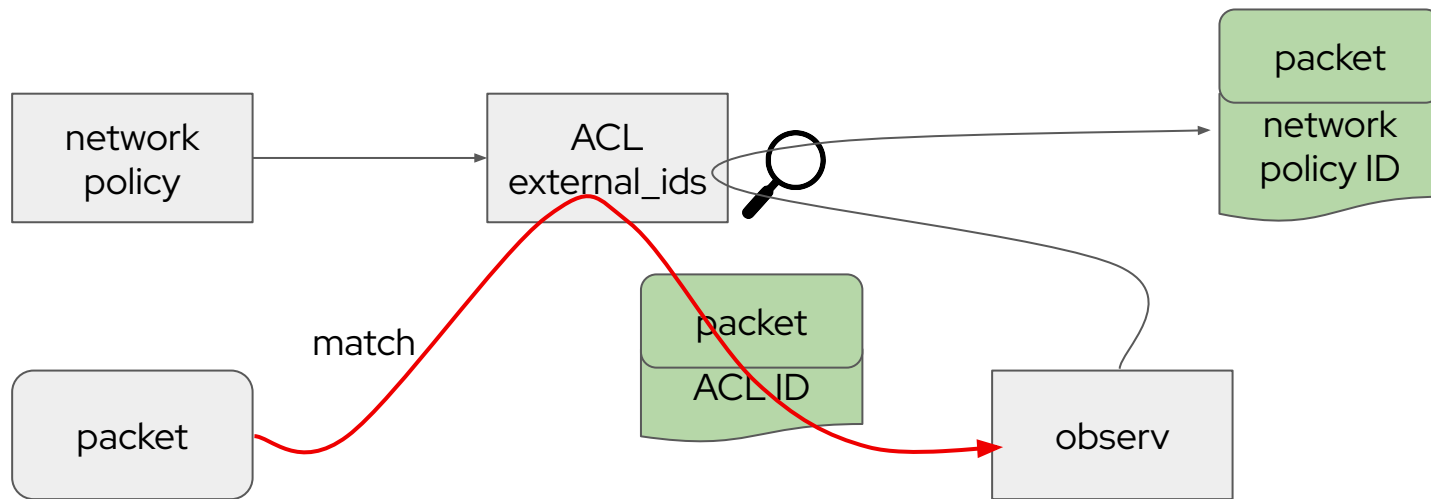


OVN-kubernetes uses ACLs to implement k8s Network Policy.

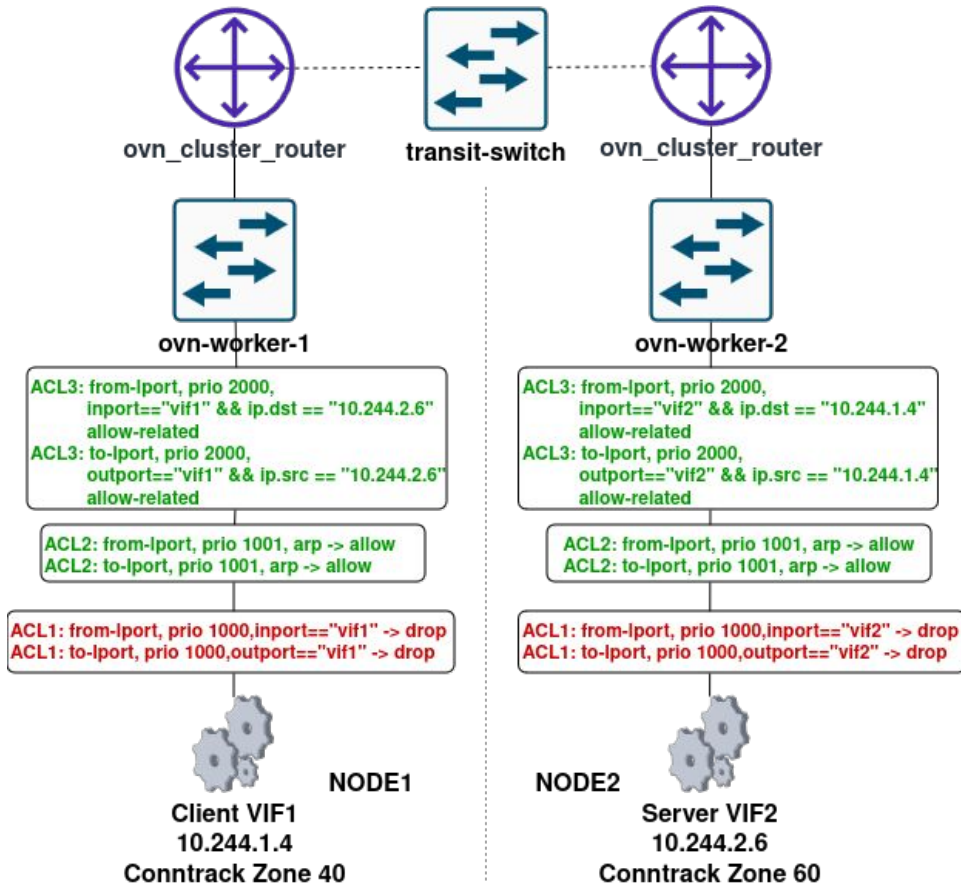




How can we do this?



# OVN ACL Sampling



- ▶ **ACL1** ("default drop")
  - lower prio (1000)
  - drops all IP traffic originating from VIFs
- ▶ **ACL2** ("allow ARPs from VIFs")
  - higher prio (1001)
- ▶ **ACL3** ("allow IP traffic between VIFs")
  - highest prio (2000)
  - only allows IP traffic between VIF1 and VIF2
  - source is selected by logical ingress port
  - destination is selected by IP
- ▶ **OVN assigned conntrack zones**
  - each VIF gets one (ovn-controller)
  - used for stateful firewalling (ACL)

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>in_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>in_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty
<b>in_acl_eval:</b> <b>match:</b> ct.new && (inport == "vif1" && ip4.dst == 10.244.2.6) (/ ACL3 *) <b>action:</b> verdict: allow, set-commit, next	zone=40, state=trk,new	empty

**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>in_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty
<b>in_acl_eval:</b> <b>match:</b> ct.new && (inport == "vif1" && ip4.dst == 10.244.2.6) /* ACL3 */ <b>action:</b> verdict: allow, set-commit, next	zone=40, state=trk,new	empty
<b>in_acl_action:</b> <b>match:</b> verdict == allow <b>action:</b> next /* allow-related */	zone=40, state=trk,new	empty
<b>in_stateful:</b> <b>match:</b> commit? <b>action:</b> commit(zone=vif1-zone), next /* allow-related */	zone=40, state=trk,new	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=new</b>

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic)

**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,est,rpl</b>	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic)

**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,est,rpl</b>	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>
<b>out_acl_hint:</b> <b>match:</b> ct.est && !ct.rel && !ct.new && !ct.inv && ct.rpl && ct_mark.blocked == 0 (/* allow-related */) <b>action:</b> verdict: allow, next	zone=40, state=trk,est,rpl	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>



**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic)

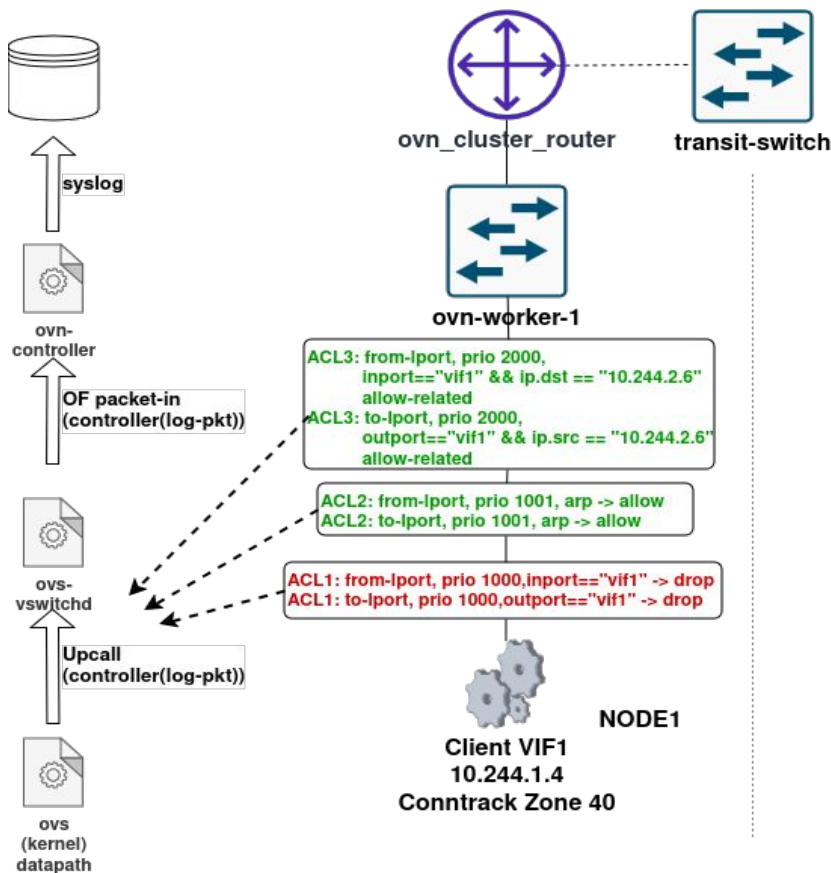
**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40 (VIF1)
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,est,rpl</b>	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>
<b>out_acl_hint:</b> <b>match:</b> ct.est && !ct.rel && !ct.new && !ct.inv && ct.rpl && ct_mark.blocked == 0 (/* allow-related */) <b>action:</b> verdict: allow, next	zone=40, state=trk,est,rpl	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>
<b>out_acl_action:</b> <b>match:</b> verdict == allow <b>action:</b> next (/* allow-related */)	zone=40, state=trk,est,rpl	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>
<b>out_stateful:</b> <b>match:</b> 1 <b>action:</b> next	zone=40, state=trk,est,rpl	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=est</b>

**OVN (pre-24.09.0)** already had features that provide visibility into ACL processing

- ACL logging
- ACL labeling

**Why not use those?**



```
$ ovn-nbctl --name ac13 --log --label 42 acl-add ...
```

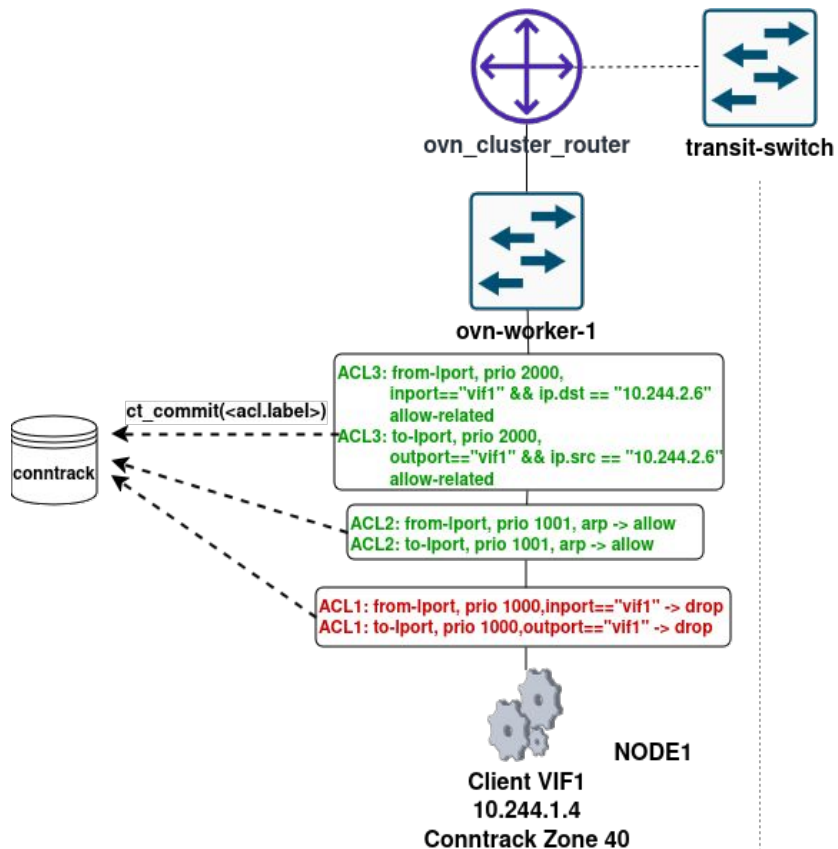
```
...
```

```
00035|acl_log(ovn_pinctrl0)|INFO|name="ac13",
verdict=allow, severity=info, direction=from-lport:
tcp,nw_src=10.244.1.4,nw_dst=10.244.2.6,tp_src=4040,
tp_dst=5201,tcp_flags=syn
```

```
00037|acl_log(ovn_pinctrl0)|INFO|name="ac13",
verdict=allow, severity=info, direction=from-lport:
tcp,nw_src=10.244.2.6,nw_dst=10.244.1.4,tp_src=5201,
tp_dst=4040,tcp_flags=syn|ack
```

## Issues

- (super) slow path all packets: upcall to ovs-vswitchd -> packet-in to ovn-controller -> syslog
- generates an additional OpenFlow rule for each ACL for reply traffic (matching on the ACL label)



```
$ ovn-nbctl --label 42 acl-add ...
```

```
...
```

```
$ ovs-appctl dpctl/dump-contrack | grep 5201
```

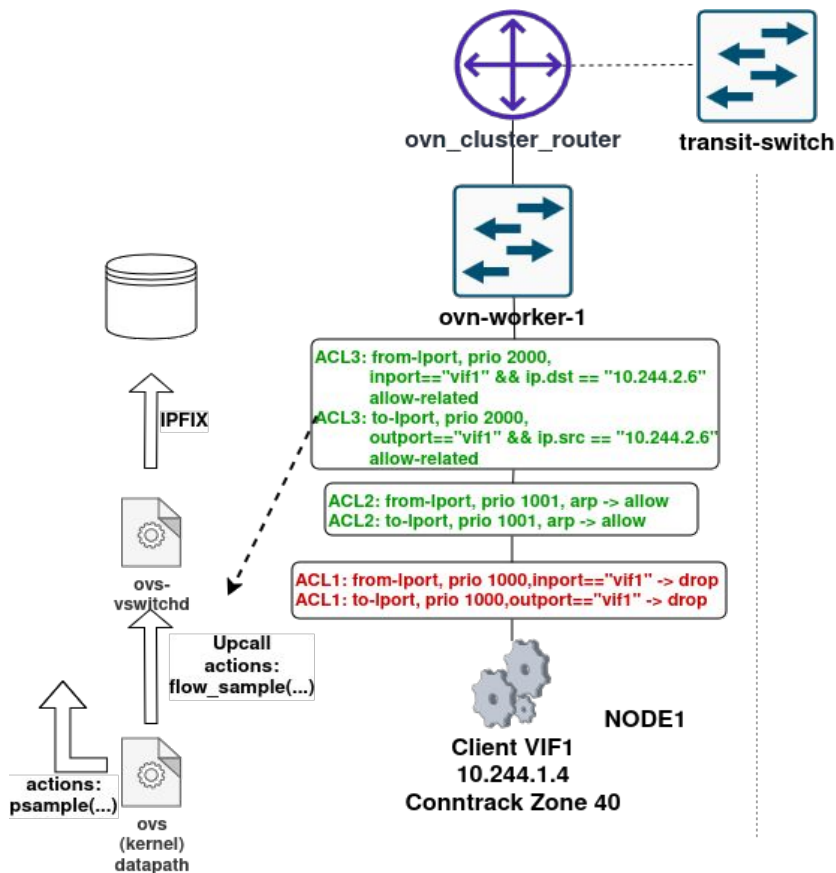
```
tcp,orig=(src=10.244.1.4,dst=10.244.2.6,sport=4040,dport=5201),reply=(src=10.244.2.6,dst=10.244.1.4,sport=5201,dport=4040),zone=2,labels=0x2a000000000000000000000000000000,proto info=(state=ESTABLISHED)
```

```
tcp,orig=(src=10.244.1.4,dst=10.244.2.6,sport=4040,dport=5201),reply=(src=10.244.2.6,dst=10.244.1.4,sport=5201,dport=4040),zone=1,labels=0x2a000000000000000000000000000000,proto info=(state=ESTABLISHED)
```

## Issues

- client applications must guess internal OVN ct\_label representation (today 32 MSB)
- no per-packet information about ACLs that were hit

**OVN 24.09 supports ACL sampling!**



```
$ ovn-nbctl list Sampling_App
```

```
id: 42, type: acl-new
```

```
id: 43, type: acl-est
```

```
$ ovn-nbctl list Sample_Collector
```

```
name: c1, probability: 100%, set_id: 100
```

```
name: c2, probability: 100%, set_id: 200
```

```
$ ovn-nbctl list Sample
```

```
uuid: s_new, collectors: [c1, c2], metadata: 1011
```

```
uuid: s_est, collectors: [c1, c2], metadata: 1012
```

```
$ ovn-nbctl list ACL
```

```
action: allow-related, direction: from-lport,
```

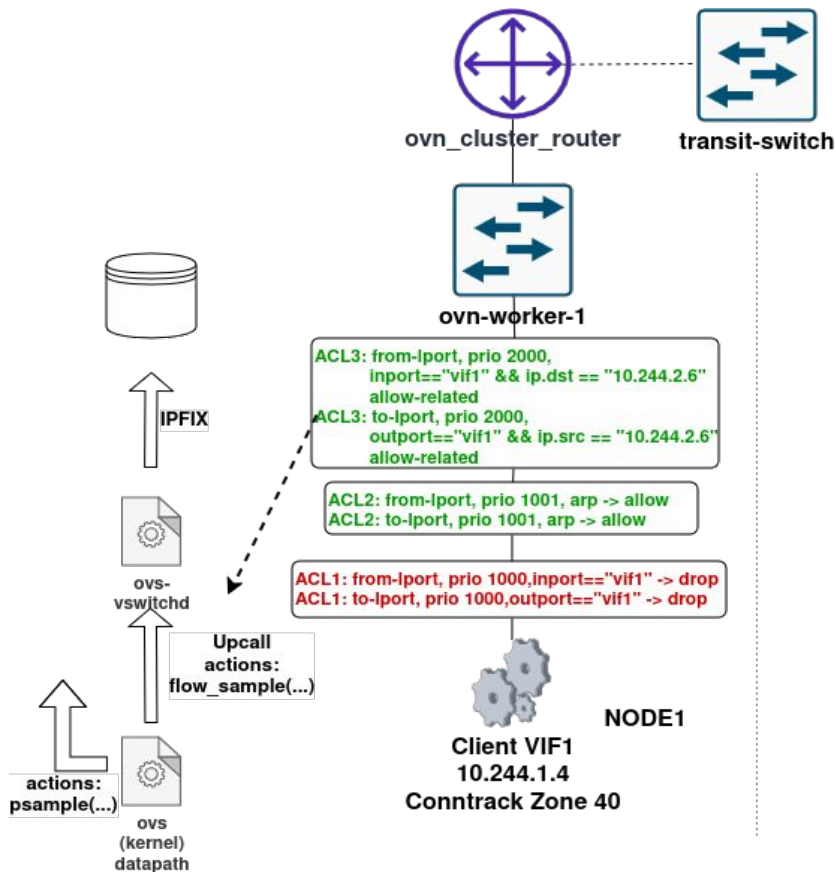
```
match: "inport == \"vif2\" && ip4.dst == 10.244.1.4",
```

```
priority: 2000, sample_new: s_new, sample_est: s_est
```

```
action: allow-related, direction: to-lport,
```

```
match: "outport == \"vif1\" && ip4.src == 10.244.2.6",
```

```
priority: 2000, sample_new: s_new, sample_est: s_est
```



```
$ ovn-nbctl list Sampling_App
```

```
id: 42, type: acl-new
```

```
id: 43, type: acl-est
```

```
$ ovn-nbctl list Sample_Collector
```

```
name: c1, probability: 100%, set_id: 100
```

```
name: c2, probability: 100%, set_id: 200
```

```
$ ovn-nbctl list Sample
```

```
uuid: s_new, collectors: [c1, c2], metadata: 1011
```

```
uuid: s_est, collectors: [c1, c2], metadata: 1012
```

```
$ ovn-nbctl list ACL
```

```
action: allow-related, direction: from-lport,
```

```
match: "inport == \"vif2\" && ip4.dst == 10.244.1.4",
```

```
priority: 2000, sample_new: s_new, sample_est: s_est
```

```
action: allow-related, direction: to-lport,
```

```
match: "outport == \"vif1\" && ip4.src == 10.244.2.6",
```

```
priority: 2000, sample_new: s_new, sample_est: s_est
```

```
$ ovs-appctl dpctl/dump-flows | grep sample
```

```
flow_sample(probability=100%,collector_set_id= 100,
```

```
obs_domain_id=0x 2b000003,obs_point_id= 1012),
```

```
flow_sample(probability= 100%,collector_set_id= 100,
```

```
obs_domain_id=0x 2b000003,obs_point_id= 1012)
```

```
flow_sample(probability= 100%,collector_set_id= 100,
```

```
obs_domain_id=0x 2a000002,obs_point_id= 1011)
```

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201, **Sampling** (new: **1011**, est: **1012**)

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40
<b>in_pre_stateful:</b> action: ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty



**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201, **Sampling** (new: 1011, est: 1012)

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40
<b>in_pre_stateful:</b> action: ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty
<b>in_acl_eval:</b> <b>match:</b> ct.new && (inport == "vif1" && ip4.dst == 10.244.2.6) (/ * ACL3 */) <b>action:</b> verdict: allow, set-commit, <b>reg3&lt;-1011, reg9&lt;-1012</b> , next	zone=40, state=trk,new	empty

**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201, **Sampling** (new: 1011, est: 1012)

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40
<b>in_pre_stateful:</b> action: ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty
<b>in_acl_eval:</b> <b>match:</b> ct.new && (inport == "vif1" && ip4.dst == 10.244.2.6) /* ACL3 */ <b>action:</b> verdict: allow, set-commit, reg3<-1011, reg9<-1012, next	zone=40, state=trk,new	empty
<b>in_acl_sample:</b> <b>match:</b> ip && ct.new && reg3 == 1011 <b>action:</b> sample(probability=100%,collector_set=200,obs_domain= <b>42</b> ,obs_point= <b>1011</b> ); sample(probability=100%,collector_set=100,obs_domain= <b>42</b> ,obs_point= <b>1011</b> );	zone=40, state=trk,new	empty

**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN 10.244.1.4:4040 -> 10.244.2.6:5201, **Sampling** (new: 1011, est: 1012)

**Outcome:** explicitly allowed by ACL3, session **committed** to conntrack

Stage + Logical Flow (ingress pipeline)	CT zone/state	CT zone 40
<b>in_pre_stateful:</b> action: ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state= <b>trk,new</b>	empty
<b>in_acl_eval:</b> <b>match:</b> ct.new && (inport == "vif1" && ip4.dst == 10.244.2.6) /* ACL3 */ <b>action:</b> verdict: allow, set-commit, reg3<-1011, reg9<-1012, next	zone=40, state=trk,new	empty
<b>in_acl_sample:</b> <b>match:</b> ip && ct.new && reg3 == 1011 <b>action:</b> sample(probability=100%,collector_set=200,obs_domain= <b>42</b> ,obs_point= <b>1011</b> ); sample(probability=100%,collector_set=100,obs_domain= <b>42</b> ,obs_point= <b>1011</b> );	zone=40, state=trk,new	empty
<b>in_acl_action:</b> match: verdict == allow <b>action:</b> next /* allow-related */	zone=40, state=trk,new	empty
<b>in_stateful:</b> match: commit? <b>action:</b> commit (zone=vif1-zone, ct_label.obs_point_id = reg9 /* 1012 */), next	zone=40, state=trk,new	10.244.1.4:4040-> 10.244.2.6:5201 <b>state=new</b>

**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic), **Sampling** (new: 1011, est: 1012)

**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state=trk,est,rpl <b>label.obs_point_id=1012</b>	state=est

**Traffic pattern:** TCP session between VIF1 -> VIF2 (**10.244.1.4:4040 -> 10.244.2.6:5201**)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic), **Sampling** (new: **1011**, est: **1012**)

**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone= <b>40</b> , state=trk,est,rpl <b>label.obs_point_id=1012</b>	<b>state=est</b>
<b>out_acl_sample:</b> <b>match:</b> ip && ct.trk && (ct.est    ct.rel) && ct.rpl && ct_label.obs_point_id == 1012 <b>action:</b> sample(probability=65535,collector_set=200,obs_domain= <b>43</b> ,obs_point= <b>1012</b> ); sample(probability=65535,collector_set=100,obs_domain= <b>43</b> ,obs_point= <b>1012</b> );	zone=40, state=trk,est,rpl	<b>state=est</b>

**Traffic pattern:** TCP session between VIF1 -> VIF2 (10.244.1.4:4040 -> 10.244.2.6:5201)

**Packet:** TCP SYN+ACK 10.244.2.6:5201 -> 10.244.1.4:4040 (reply traffic), **Sampling** (new: 1011, est: 1012)

**Outcome:** implicitly allowed by ACL3, session **found** in conntrack

Stage + Logical Flow (egress pipeline)	CT zone/state	CT zone 40
<b>out_pre_stateful:</b> <b>action:</b> ct_next(dnat) /* zone=vif1-zone */	zone=40, state=trk,est,rpl <b>label.obs_point_id=1012</b>	<b>state=est</b>
<b>out_acl_sample:</b> <b>match:</b> ip && ct.trk && (ct.est    ct.rel) && ct.rpl && ct_label.obs_point_id == 1012 <b>action:</b> sample(probability=65535,collector_set=200,obs_domain=43,obs_point=1012); sample(probability=65535,collector_set=100,obs_domain=43,obs_point=1012);	zone=40, state=trk,est,rpl	<b>state=est</b>
<b>out_acl_action:</b> <b>match:</b> verdict == allow <b>action:</b> next /* allow-related */	zone=40, state=trk,est,rpl	<b>state=est</b>
<b>out_stateful:</b> <b>match:</b> 1 <b>action:</b> next	zone=40, state=trk,est,rpl	<b>state=est</b>

## The good:

- ▶ can provide per packet ACL information
- ▶ can distinguish between packets that create a **new** connection OR are part of an **established** connection
- ▶ datapath agnostic: works fine with both **IPFIX** and **psample** collectors

## The not so good:

- ▶ additional flows needed (one per *ACL - for sample\_est -*) needed in the *reply direction*  
=> O(number-of-ACLs) extra flows => **control plane cost!**

**OVS v3.4.0** includes [1aa9e137fe36 \("ofp-actions: Load data from fields in sample action."\)](#)

OVN rewrites the logical flows in the reply direction (for *ACL Sample configurations with exactly one collector*):

► **From one flow per ACL:**

out\_acl\_sample:

match: ip && ct.trk && (ct.est || ct.rel) && ct.rpl && ct\_label.obs\_point\_id == **1012**

action: sample(probability=65535,collector\_set=100,obs\_domain=43,obs\_point=**1012**);

► **To exactly two flows (regardless of number of ACLs):**

out\_acl\_sample:

match=ip && ct.trk && (ct.est || ct.rel) && !ct.rpl && ct\_mark.obs\_collector\_id == 1

action=sample(probability=100%,collector\_set=100,obs\_domain=43,obs\_point=**ct\_label.obs\_point\_id**);

out\_acl\_sample:

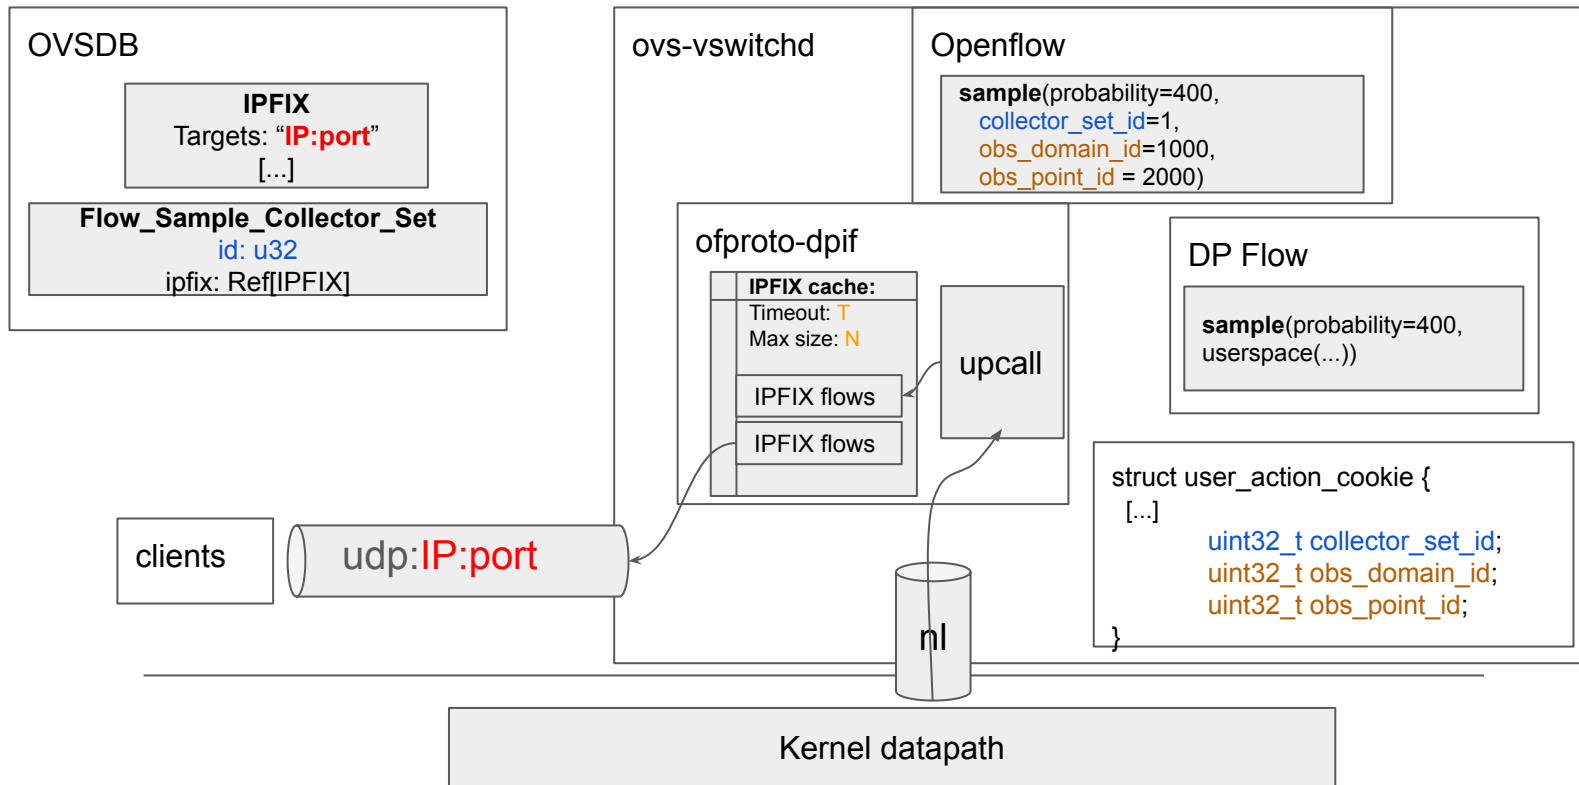
match=ip && ct.trk && (ct.est || ct.rel) && ct.rpl && ct\_mark.obs\_collector\_id == 1

action=sample(probability=65535,collector\_set=100,obs\_domain=43,obs\_point=**ct\_label.obs\_point\_id**);



# OVS (p)sampling

## Existing sampling infrastructure: IPFIX



## Main problems

- ▶ Each sample goes through ovs-vswitchd: overload
- ▶ Samples share netlink socket with upcalls: packet loss!
- ▶ IPFIX cache accessed *handler* threads: contention
- ▶ IPFIX socket sent from *main* or *handler* threads

## Goal

- ▶ Send the sample to the end user directly (bypassing ovs-vswitchd)
- ▶ Have the sample contain metadata extracted from OpenFlow

## Introducing psample

- ✓ psample exists from v4.11, used by tc (see [tc\\_sample\(8\)](#))
- ✓ Creates a netlink multicast group where samples are sent
- ✓ A (u32) "group id" identifies "source" of the sample
- ✓ Includes sampling rate
  
- ✗ It does not have a cookie or any user-defined metadata
- ✗ The packet is copied even if there are no listeners

## Introducing psample

- ✓ psample exists from v4.11, used by tc (see [tc\\_sample\(8\)](#))
- ✓ Creates a netlink multicast group where samples are sent
- ✓ A (u32) "group id" identifies "source" of the sample
- ✓ Includes sampling rate

**added!**

**optimized!**

It does not have a cookie or any user-defined metadata

The packet is copied even if there are no listeners

## OVS datapath psample action

```
psample(group=42, cookie=0x1234)
```

- ▶ It integrates with other datapath actions. E.g:
  - `sample(sample=10%, actions(psample(cookie=0x123)))`
  - `trunc(20), psample()`
- ▶ Only implemented in kernel datapath \*

# OpenFlow sample action optimization

```
sample(obs_domain_id=42,obs_point_id = 2000)
```



## OpenFlow sample action optimization

```
ct_label = 1000, actions=sample(obs_domain_id=42, obs_point_id=1000)  
ct_label = 2000, actions=sample(obs_domain_id=42, obs_point_id=2000)  
ct_label = 3000, actions=sample(obs_domain_id=42, obs_point_id=3000)  
...
```

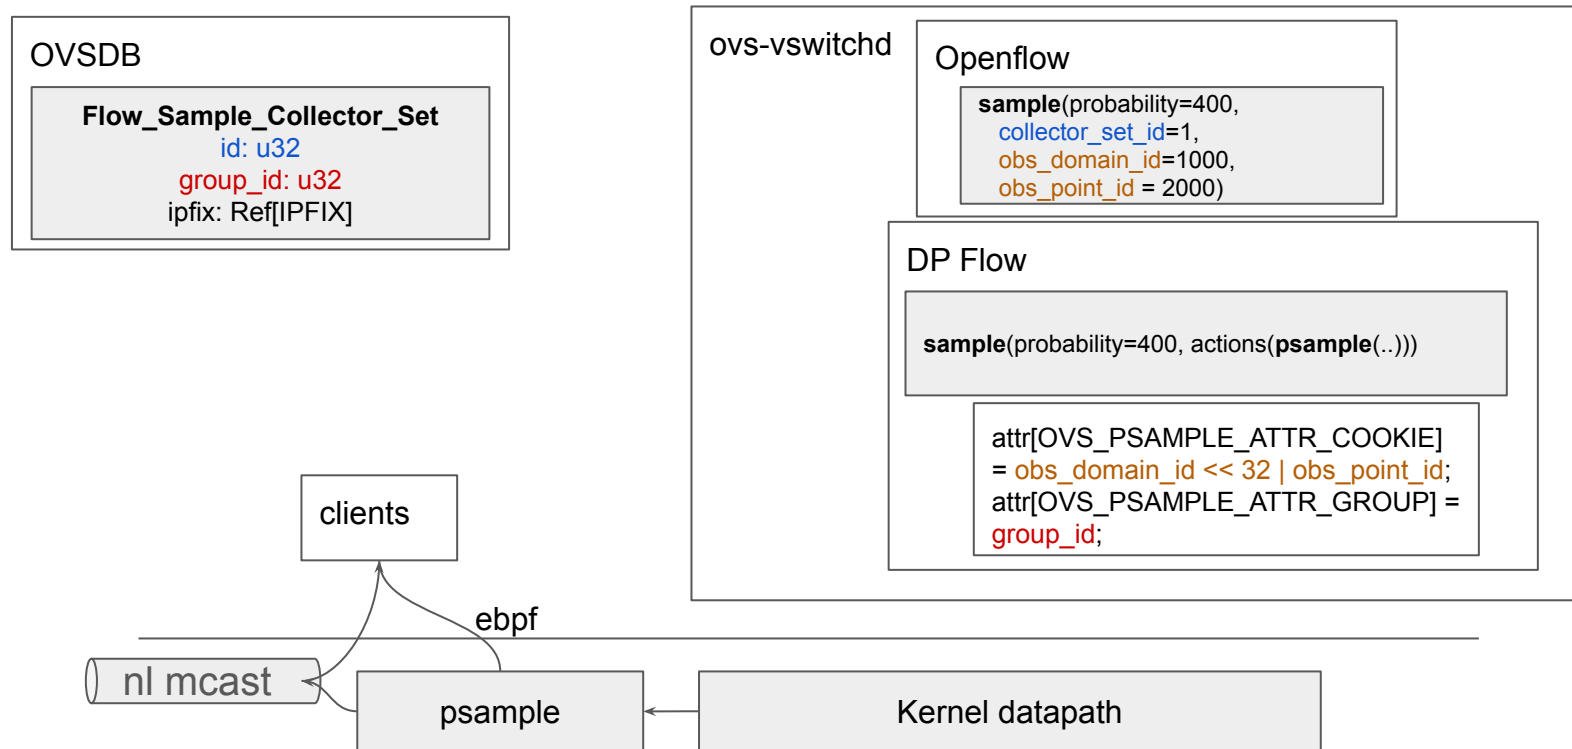
## OpenFlow sample action optimization

```
ct_label = 1000, actions=sample(obs_domain_id=42, obs_point_id=1000)  
ct_label = 2000, actions=sample(obs_domain_id=42, obs_point_id=2000)  
ct_label = 3000, actions=sample(obs_domain_id=42, obs_point_id=3000)  
...
```



```
actions=sample(obs_domain_id=42, obs_point_id=NXM_NX_CT_LABEL[32..63])
```

## New sampling infrastructure: psample



## Future work

- ▶ Finalize and publish performance tests. Stay tuned!
- ▶ Explore a userspace implementation
  - USDT probes? Userspace ebpf VM?
  - Static (hardcoded) behavior?
  - Ideas and requirements are welcome!

# Demo

## Network Traffic

Query options ▾

Quick filters ▾

Source Namespace ▾

= ▾

E.g: netobserv ▾



[Hide filters](#)

Time range

Last 5 minutes ▾

Refresh Interval

Refresh off ▾



[More options](#)

Specify a single kubernetes name.

[Learn more](#)

Source Namespace

iperf



[More options](#)

Overview

**Traffic flows**

Topology

[Show histogram](#)

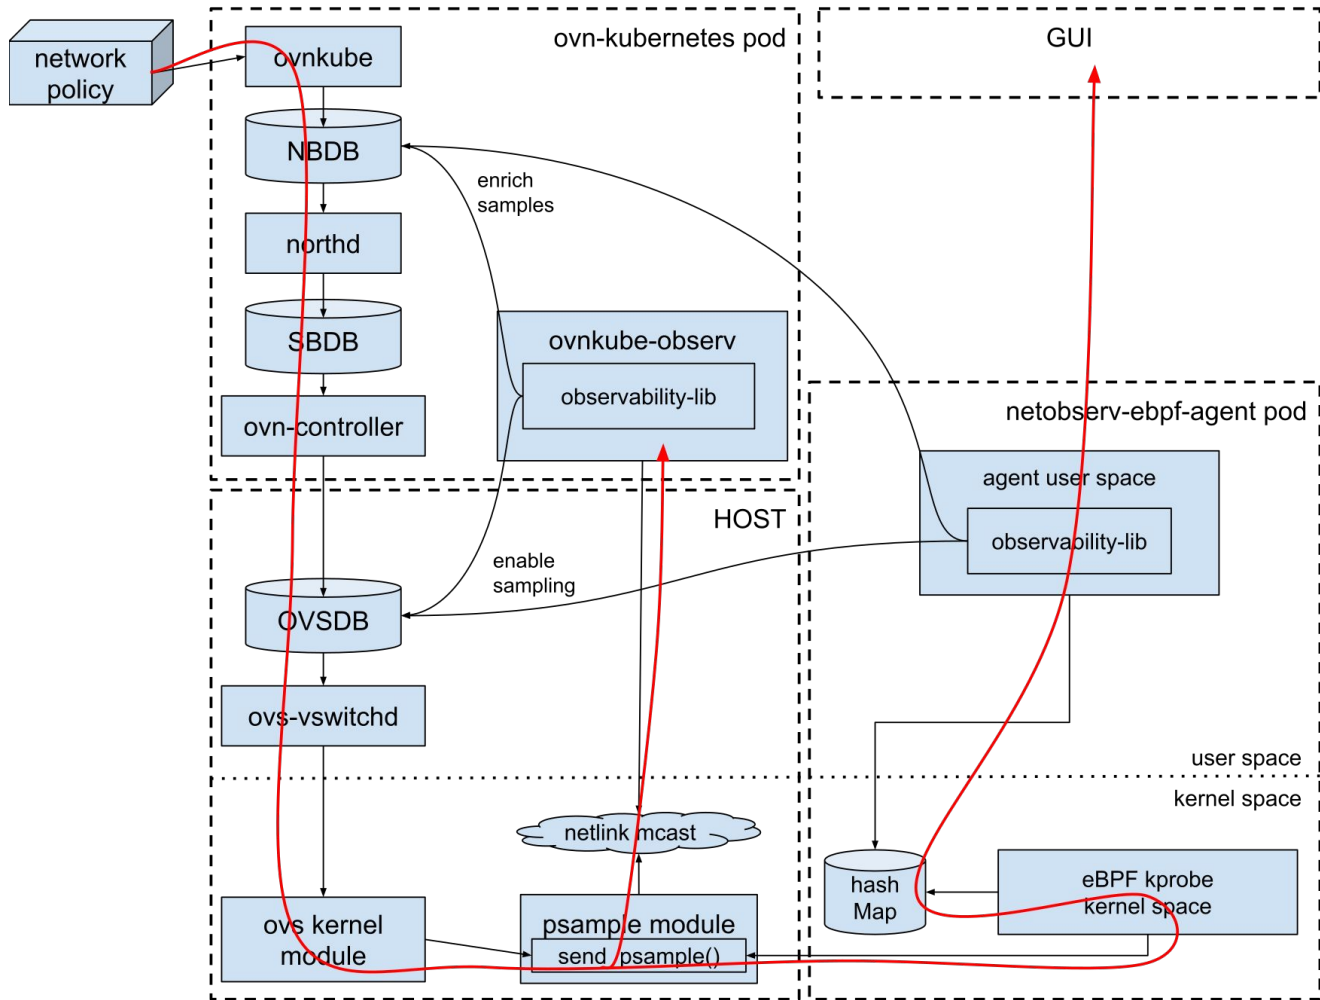
[Hide advanced options](#)

Display options ▾

[More options](#)

Source		Destination			L3 Layer	Packets	Network Events
Name		Name	Namespace	Port	Protocol		
iperf3-clients-jlbp5		iperf3-server-deployment-55cc588d8b-mtmnj	iperf	5201	TCP	1	Allowed by network policy iperf:iperf3-server-access-egress, direction Egress
iperf3-clients-jlbp5		iperf3-server-deployment-55cc588d8b-mtmnj	iperf	5201	TCP	1	Dropped by network policies isolation in namespace iperf, direction Ingress







# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)